

intro_numpy

January 10, 2021

```
[2]: # The following is to know when this notebook has been run and with which
      ↪python version.
import time, sys
print(time.ctime())
print(sys.version.split('|')[0])
```

Mon Oct 12 19:08:54 2020

3.7.6 (default, Jan 8 2020, 13:42:34)

[Clang 4.0.1 (tags/RELEASE_401/final)]

1 B Numpy

This is part of the Python lecture given by Christophe Morisset at IA-UNAM.

1.0.1 Import numpy first

```
[3]: # You need first to import the numpy library (must be installed on your
      ↪computer ;- )
      # As it will be widely used, better to give it a nickname, or an alias.
      ↪Traditionnaly, it's "np":
import numpy as np
```

```
[4]: print(np.__version__)
```

1.18.1

1.0.2 Tutorials

<http://nbviewer.ipython.org/github/jrjohansson/scientific-python-lectures/blob/master/Lecture-2-Numpy.ipynb> AND <http://nbviewer.ipython.org/gist/rpmuller/5920182> AND http://www.astro.washington.edu/users/vanderplas/Astr599/notebooks/11_EfficientNumpy

1.0.3 The ARRAY class

Create an array

```
[5]: # Easy to create a numpy array (the basic class) from a list
l = [1,2,3,4,5,6]
print(l)
```

```
a = np.array([1,2,3,4,5,6])
print(a)
print(type(a))
```

```
[1, 2, 3, 4, 5, 6]
[1 2 3 4 5 6]
<class 'numpy.ndarray'>
```

```
[6]: # Works with tuples also:
b = np.array((1,2,3))
print(b)
```

```
[1 2 3]
```

Numpy arrays are efficiently connected to the computer:

```
[7]: L = range(1000)
%timeit L2 = [i**2 for i in L] # Notice the use of timeit, a magic function.
↳ (starts with %)
A = np.arange(1000)
%timeit A2 = A**2
```

```
293 µs ± 19.5 µs per loop (mean ± std. dev. of 7 runs, 1000 loops each)
842 ns ± 7.13 ns per loop (mean ± std. dev. of 7 runs, 1000000 loops each)
```

```
[8]: L = [1, 2, 3, 4]
a = np.array(L)
print(a.dtype)
print(a)
```

```
int64
[1 2 3 4]
```

```
[9]: L = [1, 2, 3, 4.]
a = np.array(L)
print(a.dtype)
print(a)
```

```
float64
[1. 2. 3. 4.]
```

```
[10]: L = [1, 2, 3, 4.5, 'a']
a = np.array(L)
print(L) # Different types can coexist in a python list
print(a.dtype)
print(a) # NOT in a numpy array. The array is re-typed to the highest type,
↳ here string.
```

```
[1, 2, 3, 4.5, 'a']
<U32
['1' '2' '3' '4.5' 'a']
```

Once the type of an array is defined, one can insert values of type that can be transformed to the type of the array

```
[11]: a = np.array([1, 2, 3, 4, 5, 6])
      print(a)
      a[4] = 2.56 # will be transformed to int(2.56)
      print(a)
      a[3] = '20' # will be transformed to int('20')
      print(a)
```

```
[1 2 3 4 5 6]
[1 2 3 4 2 6]
[ 1  2  3 20  2  6]
```

```
[12]: a[2] = '3.2'
```

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-12-81ec24b7f2d9> in <module>
----> 1 a[2] = '3.2'

ValueError: invalid literal for int() with base 10: '3.2'
```

```
[15]: a[2] = 'tralala'
```

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-15-2774d8249152> in <module>
----> 1 a[2] = 'tralala'

ValueError: invalid literal for int() with base 10: 'tralala'
```

1D, 2D, 3D, ...

```
[16]: a = np.array([1,2,3,4,5,6])
      b = np.array([[1,2],[1,4],[4,7]])
      c = np.array([[[1],[2]], [[3],[4]]])
      print(a.shape, b.shape, c.shape)
      print(b)
      print(b[0]) # no error
```

```
(6,) (3, 2) (2, 2, 1)
[[1 2]
```

```
[1 4]
[4 7]]
[1 2]
```

```
[17]: print(len(a), len(b), len(c)) # size of the first dimension
```

```
6 3 2
```

```
[18]: b.size
```

```
[18]: 6
```

```
[19]: print(a.ndim, b.ndim, c.ndim)
```

```
1 2 3
```

```
[20]: a = np.array([1,2,3,4,5,6])
print('mean: {0}, max: {1}, shape: {2}, median: {3}'.format(a.mean(), a.max(),
↳a.shape, np.median(a)))
```

```
mean: 3.5, max: 6, shape: (6,), median: 3.5
```

mean and max are methods (functions) of the array class, they need (). shape is an attribute (like a variable).

```
[21]: print(a.mean) # this is printing information about the function, NOT the result
↳of the function!
```

```
<built-in method mean of numpy.ndarray object at 0x7fd0f9fb3990>
```

```
[22]: mm = a.mean # We assign to mm the function. Then we can call it directly, but
↳still need for the ():
print(mm())
```

```
3.5
```

```
[23]: print(b)
print(b.mean()) # mean over the whole array
print(b.mean(axis=0)) # mean over the first axis (columns)
print(b.mean(1)) # mean over the rows
print(np.mean(b, axis=1))
```

```
[[1 2]
 [1 4]
 [4 7]]
3.1666666666666665
[2.          4.33333333]
[1.5 2.5 5.5]
[1.5 2.5 5.5]
```

```
[24]: np.mean?
```

Signature: `np.mean(a, axis=None, dtype=None, out=None, keepdims=<no value>)`

Docstring:

Compute the arithmetic mean along the specified axis.

Returns the average of the array elements. The average is taken over the flattened array by default, otherwise over the specified axis. ``float64`` intermediate and return values are used for integer inputs.

Parameters

`a` : array_like

Array containing numbers whose mean is desired. If ``a`` is not an array, a conversion is attempted.

`axis` : None or int or tuple of ints, optional

Axis or axes along which the means are computed. The default is to compute the mean of the flattened array.

.. versionadded:: 1.7.0

If this is a tuple of ints, a mean is performed over multiple axes, instead of a single axis or all the axes as before.

`dtype` : data-type, optional

Type to use in computing the mean. For integer inputs, the default is ``float64``; for floating point inputs, it is the same as the input dtype.

`out` : ndarray, optional

Alternate output array in which to place the result. The default is ``None``; if provided, it must have the same shape as the expected output, but the type will be cast if necessary. See ``ufuncs-output-type`` for more details.

`keepdims` : bool, optional

If this is set to True, the axes which are reduced are left in the result as dimensions with size one. With this option, the result will broadcast correctly against the input array.

If the default value is passed, then ``keepdims`` will not be passed through to the ``mean`` method of sub-classes of ``ndarray``, however any non-default value will be. If the sub-class' method does not implement ``keepdims`` any exceptions will be raised.

Returns

`m` : ndarray, see dtype parameter above

If ``out=None``, returns a new array containing the mean values, otherwise a reference to the output array is returned.

See Also

average : Weighted average
std, var, nanmean, nanstd, nanvar

Notes

The arithmetic mean is the sum of the elements along the axis divided by the number of elements.

Note that for floating-point input, the mean is computed using the same precision the input has. Depending on the input data, this can cause the results to be inaccurate, especially for ``float32`` (see example below). Specifying a higher-precision accumulator using the ``dtype`` keyword can alleviate this issue.

By default, ``float16`` results are computed using ``float32`` intermediates for extra precision.

Examples

```
>>> a = np.array([[1, 2], [3, 4]])
>>> np.mean(a)
2.5
>>> np.mean(a, axis=0)
array([2., 3.])
>>> np.mean(a, axis=1)
array([1.5, 3.5])
```

In single precision, ``mean`` can be inaccurate:

```
>>> a = np.zeros((2, 512*512), dtype=np.float32)
>>> a[0, :] = 1.0
>>> a[1, :] = 0.1
>>> np.mean(a)
0.54999924
```

Computing the mean in float64 is more accurate:

```
>>> np.mean(a, dtype=np.float64)
0.550000000074505806 # may vary
File:      ~/anaconda3/lib/python3.7/site-packages/numpy/core/fromnumeric.py
Type:      function
```

Creating arrays from scratch

```
[25]: print(np.arange(10))
```

```
[0 1 2 3 4 5 6 7 8 9]
```

```
[26]: print(np.linspace(0, 1, 10)) # start, stop (included), number of points
print('-----')
print(np.linspace(0, 1, 11)) # start, stop (included), number of points
print('-----')
print(np.linspace(0, 1, 10, endpoint=False)) # Not including the stop point
```

```
[0.          0.11111111 0.22222222 0.33333333 0.44444444 0.55555556
 0.66666667 0.77777778 0.88888889 1.          ]
```

```
-----
```

```
[0.  0.1 0.2 0.3 0.4 0.5 0.6 0.7 0.8 0.9 1. ]
```

```
-----
```

```
[0.  0.1 0.2 0.3 0.4 0.5 0.6 0.7 0.8 0.9]
```

```
[27]: print(np.logspace(0, 2, 10)) # from 10**start to 10**stop, with 10 values
```

```
[ 1.          1.66810054  2.7825594   4.64158883  7.74263683
 12.91549665  21.5443469   35.93813664  59.94842503 100.          ]
```

```
[28]: print(np.zeros(2)) # Filled with 0.0
print('-----')
print(np.zeros((2,3))) # a 2D array, also filled with 0.0
print('-----')
print(np.ones_like(b)) # This is very usefull: using an already created array
    ↳ (or list or tuple) as example for the shape of the new one.
print('-----')
print(b)
print(np.zeros_like(b, dtype=float) + 3) # Can define a value to fille the
    ↳ array when creating it. Or latter...
print(np.ones_like(b, dtype=float) * 3) # Can define a value to fille the array
    ↳ when creating it. Or latter...
```

```
[0. 0.]
```

```
-----
```

```
[[0. 0. 0.]
 [0. 0. 0.]]
```

```
-----
```

```
[[1 1]
 [1 1]
 [1 1]]
```

```
-----
```

```
[[1 2]
 [1 4]
 [4 7]]
```

```
[[3. 3.]
 [3. 3.]
 [3. 3.]]
[[3. 3.]
 [3. 3.]
 [3. 3.]]
```

```
[29]: new_a = a.reshape((3,2)) # This does NOT change the shape of a
print(a)
print('-----')
print(new_a)
```

```
[1 2 3 4 5 6]
-----
[[1 2]
 [3 4]
 [5 6]]
```

```
[30]: print(new_a.ravel())
print(new_a.reshape(new_a.size))
```

```
[1 2 3 4 5 6]
[1 2 3 4 5 6]
```

```
[31]: # create 2 2D arrays (coordinates matrices), one describing how x varies, the
      ↪ other for y.
x, y = np.mgrid[0:5, 0:10] # This is not a function!!! notice the []
print(x.shape)
print(x)
print('-----')
print(y.shape)
print(y)
```

```
(5, 10)
[[0 0 0 0 0 0 0 0 0 0]
 [1 1 1 1 1 1 1 1 1 1]
 [2 2 2 2 2 2 2 2 2 2]
 [3 3 3 3 3 3 3 3 3 3]
 [4 4 4 4 4 4 4 4 4 4]]
```

```
-----
(5, 10)
[[0 1 2 3 4 5 6 7 8 9]
 [0 1 2 3 4 5 6 7 8 9]
 [0 1 2 3 4 5 6 7 8 9]
 [0 1 2 3 4 5 6 7 8 9]
 [0 1 2 3 4 5 6 7 8 9]]
```



```
[32]: # coordinates matrices using user-defined x- and y-vectors
x, y = np.meshgrid([1,2,4,7], [0.1, 0.2, 0.3])
print(x)
print('-----')
print(y)
```

```
[[1 2 4 7]
 [1 2 4 7]
 [1 2 4 7]]
-----
[[0.1 0.1 0.1 0.1]
 [0.2 0.2 0.2 0.2]
 [0.3 0.3 0.3 0.3]]
```

```
[33]: x, y = np.meshgrid([1,2,4,7], [0.1, 0.2, 0.3], indexing='ij') # the other order.
      ↪ ...
print(x)
print('-----')
print(y)
```

```
[[1 1 1]
 [2 2 2]
 [4 4 4]
 [7 7 7]]
-----
[[0.1 0.2 0.3]
 [0.1 0.2 0.3]
 [0.1 0.2 0.3]
 [0.1 0.2 0.3]]
```

WARNING arrays share memory

```
[34]: b = a.reshape((3,2))
print(a.shape, b.shape)
```

```
(6,) (3, 2)
```

```
[35]: b[1,1] = 100 # modify a value in the array
print(b)
```

```
[[ 1  2]
 [ 3 100]
 [ 5  6]]
```

```
[36]: print(a) # !!! a and b are sharing the same place in the memory, they are ↪
      ↪ pointing to the same values.
```

```
[ 1  2  3 100  5  6]
```

```
[37]: b[1,1], a[3] # same value
```

```
[37]: (100, 100)
```

```
[38]: a is b # a and b are different
```

```
[38]: False
```

```
[39]: print(b[1,1] == a[3])  
print(b[1,1] is a[3]) # Even if the values are the same, the "is" does not tell  
→ it.
```

```
True
```

```
False
```

```
[40]: c = a.reshape((2,3)).copy() # This is the solution.
```

```
[41]: print(a)  
print('-----')  
print(c)
```

```
[ 1  2  3 100  5  6]
```

```
-----
```

```
[[ 1  2  3]  
 [100  5  6]]
```

```
[42]: c[0,0] = 8888  
print(a)  
print('-----')  
print(c)
```

```
[ 1  2  3 100  5  6]
```

```
-----
```

```
[[8888  2  3]  
 [ 100  5  6]]
```

1.0.4 Random

```
[43]: ran_uniform = np.random.rand(5) # between 0 and 1  
ran_normal = np.random.randn(5) # Gaussian mean 0 variance 1  
print(ran_uniform)  
print('-----')  
print(ran_normal)  
print('-----')  
ran_normal_2D = np.random.randn(5,5) # Gaussian mean 0 variance 1  
print(ran_normal_2D)
```

```
[0.00493762 0.56611721 0.73756545 0.87137075 0.97132855]
```

```
-----  
[-0.51035597 -0.22744195 -0.44674488 -2.11004404 0.03959021]  
-----
```

```
[-0.14487643 1.41711047 -0.19250961 2.72138627 -1.32575917]  
[ 1.308128 0.33684976 0.85683118 0.67287322 -0.13687007]  
[-2.37974623 -0.9063964 -1.1233723 1.42649559 -0.16433894]  
[-0.43420222 -0.77281828 -1.57142812 0.09280614 -1.19114424]  
[ 0.29052205 -0.19884749 0.1487806 1.48351285 1.95064135]]
```

```
[44]: np.random.seed(1)  
print(np.random.rand(5))  
#np.random.seed(1)  
print(np.random.rand(5))
```

```
[4.17022005e-01 7.20324493e-01 1.14374817e-04 3.02332573e-01  
1.46755891e-01]  
[0.09233859 0.18626021 0.34556073 0.39676747 0.53881673]
```

1.0.5 Timing on 2D array

```
[45]: N = 100  
A = np.random.rand(N, N)  
B = np.zeros_like(A)
```

```
[46]: %%timeit  
for i in range(N):  
    for j in range(N):  
        B[i,j] = A[i,j]
```

2.38 ms \pm 131 μ s per loop (mean \pm std. dev. of 7 runs, 100 loops each)

```
[47]: %%timeit  
B = A # very faster ! It does NOT copy...
```

19.3 ns \pm 0.222 ns per loop (mean \pm std. dev. of 7 runs, 10000000 loops each)

```
[48]: %%timeit  
B = (A.copy()) # Takes more time
```

2.61 μ s \pm 210 ns per loop (mean \pm std. dev. of 7 runs, 100000 loops each)

```
[49]: %%timeit  
for i in range(N):  
    for j in range(N):  
        B[i,j] = A[i,j]**2
```

4.85 ms \pm 39.7 μ s per loop (mean \pm std. dev. of 7 runs, 100 loops each)

```
[50]: %%timeit
      B = A**2 # very faster ! Does a copy
```

5.69 μ s \pm 193 ns per loop (mean \pm std. dev. of 7 runs, 100000 loops each)

```
[51]: %timeit B = (A.copy())**2 # Takes a little bit more time
```

8.75 μ s \pm 51 ns per loop (mean \pm std. dev. of 7 runs, 100000 loops each)

1.0.6 Slicing

```
[52]: a = np.arange(10)
      print(a)
      print(a[1:8:3])
```

```
[0 1 2 3 4 5 6 7 8 9]
[1 4 7]
```

```
[53]: print(a[:7])
```

```
[0 1 2 3 4 5 6]
```

```
[54]: print(a[4:])
```

```
[4 5 6 7 8 9]
```

```
[55]: print(a[:2])
      print(a[:2][2])
```

```
[0 2 4 6 8]
4
```

```
[56]: # Revert the array:
      print(a[::-1])
```

```
[9 8 7 6 5 4 3 2 1 0]
```

Assignment

```
[57]: a[5:] = 999
      print(a)
```

```
[ 0  1  2  3  4 999 999 999 999 999]
```

```
[58]: a[5:] = a[4::-1]
      print(a)
```

```
[0 1 2 3 4 4 3 2 1 0]
```

```
[59]: print(a)
      b = a[:, np.newaxis] # create a new empty dimension
      print(b)
      print(a.shape, b.shape)
      c = a[np.newaxis, :]
      print(c, c.shape)
```

```
[0 1 2 3 4 4 3 2 1 0]
[[0]
 [1]
 [2]
 [3]
 [4]
 [4]
 [3]
 [2]
 [1]
 [0]]
(10,) (10, 1)
[[0 1 2 3 4 4 3 2 1 0]] (1, 10)
```

```
[60]: b * c # Cross product, see below (broadcasting)
```

```
[60]: array([[ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0],
             [ 0,  1,  2,  3,  4,  4,  3,  2,  1,  0],
             [ 0,  2,  4,  6,  8,  8,  6,  4,  2,  0],
             [ 0,  3,  6,  9, 12, 12,  9,  6,  3,  0],
             [ 0,  4,  8, 12, 16, 16, 12,  8,  4,  0],
             [ 0,  4,  8, 12, 16, 16, 12,  8,  4,  0],
             [ 0,  3,  6,  9, 12, 12,  9,  6,  3,  0],
             [ 0,  2,  4,  6,  8,  8,  6,  4,  2,  0],
             [ 0,  1,  2,  3,  4,  4,  3,  2,  1,  0],
             [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0]])
```

Using an array

```
[61]: print(a)
      a[[2,4,6]] = -999
      print(a)
```

```
[0 1 2 3 4 4 3 2 1 0]
[  0  1 -999  3 -999  4 -999  2  1  0]
```

```
[62]: # a = 1 would turn a to be 1, but if we want to assign 1 to every value in a
      ↪ one must do:
      a[:] = 1
      print(a)
```

```
[1 1 1 1 1 1 1 1 1 1]
```

1.0.7 Using masks

```
[63]: a = np.random.randint(0, 100, 20) # min, max, N  
      print(a)
```

```
[78 46 26 63 86  2 96 45 13 67 37 36 54 63 65 58 49 48 59 26]
```

```
[64]: a < 50
```

```
[64]: array([False,  True,  True, False, False,  True, False,  True,  True,  
          False,  True,  True, False, False, False, False,  True,  True,  
          False,  True])
```

```
[65]: mask = (a < 50)
```

```
[66]: mask.sum()
```

```
[66]: 10
```

```
[67]: a[mask]
```

```
[67]: array([46, 26,  2, 45, 13, 37, 36, 49, 48, 26])
```

```
[68]: b = a.copy() # do NOT use b = a  
      b[mask] = 50 #  
      print(a)  
      print(b)
```

```
[78 46 26 63 86  2 96 45 13 67 37 36 54 63 65 58 49 48 59 26]
```

```
[78 50 50 63 86 50 96 50 50 67 50 50 54 63 65 58 50 50 59 50]
```

```
[69]: b = a.copy()  
      b[b <= 50] = -1 # shortest way. Not matter if not even one element fit the test  
      print(a)  
      print(b)
```

```
[78 46 26 63 86  2 96 45 13 67 37 36 54 63 65 58 49 48 59 26]
```

```
[78 -1 -1 63 86 -1 96 -1 -1 67 -1 -1 54 63 65 58 -1 -1 59 -1]
```

```
[70]: print(a[mask])  
      print(a[~mask]) # complementary
```

```
[46 26  2 45 13 37 36 49 48 26]
```

```
[78 63 86 96 67 54 63 65 58 59]
```

```
[71]: mask
```

```
[71]: array([False,  True,  True, False, False,  True, False,  True,  True,
          False,  True,  True, False, False, False, False,  True,  True,
          False,  True])
```

```
[72]: mask = np.zeros_like(a, dtype=bool)
      print(mask)
```

```
[False False False False False False False False False False False False
 False False False False False False False False]
```

```
[73]: mask[[2,3,4]] = True
```

```
[74]: mask
```

```
[74]: array([False, False,  True,  True,  True, False, False, False, False,
          False, False, False, False, False, False, False, False, False,
          False, False])
```

```
[75]: a[mask]
```

```
[75]: array([26, 63, 86])
```

```
[76]: a[mask].sum()
```

```
[76]: 175
```

```
[77]: a[mask].mean()
```

```
[77]: 58.333333333333336
```

combining masks

```
[78]: print(a)
      mask_low = a > 30
      mask_high = a < 70
      print('-----')
      print(a[mask_low & mask_high]) # both conditions are filled
      print('-----')
      print(a[~mask_low | ~mask_high]) # complementary, using the | for OR
```

```
[78 46 26 63 86  2 96 45 13 67 37 36 54 63 65 58 49 48 59 26]
```

```
-----
```

```
[46 63 45 67 37 36 54 63 65 58 49 48 59]
```

```
-----
```

```
[78 26 86  2 96 13 26]
```

the where function

```
[79]: tt = np.where(a > 30)
      print(a)
      print(tt) # tt is a tuple of arrays, one for each dimension of the condition,
                # containing the indices where the condition is filled in that dimension.
```

```
[78 46 26 63 86  2 96 45 13 67 37 36 54 63 65 58 49 48 59 26]
(array([ 0,  1,  3,  4,  6,  7,  9, 10, 11, 12, 13, 14, 15, 16, 17, 18]),)
```

```
[80]: (a > 30).nonzero() # "where" is the same than condition.nonzero().
```

```
[80]: (array([ 0,  1,  3,  4,  6,  7,  9, 10, 11, 12, 13, 14, 15, 16, 17, 18]),)
```

```
[81]: # the indices where the condition is filled are in the first element of the
      ↪ tuple
```

```
[82]: tt[0]
```

```
[82]: array([ 0,  1,  3,  4,  6,  7,  9, 10, 11, 12, 13, 14, 15, 16, 17, 18])
```

```
[83]: # faster once you know that the condition is 1D
      tt = np.where(a > 30)[0]
```

```
[84]: tt # the array containing the indices where the condition is filled
```

```
[84]: array([ 0,  1,  3,  4,  6,  7,  9, 10, 11, 12, 13, 14, 15, 16, 17, 18])
```

```
[85]: a[tt] # the values where the condition is filled
```

```
[85]: array([78, 46, 63, 86, 96, 45, 67, 37, 36, 54, 63, 65, 58, 49, 48, 59])
```

```
[86]: # The where function can take 3 arguments.
      b = np.where(a < 50, np.nan, a)
      print(a)
      print(b)
      print(np.isfinite(b))
```

```
[78 46 26 63 86  2 96 45 13 67 37 36 54 63 65 58 49 48 59 26]
[78. nan nan 63. 86. nan 96. nan nan 67. nan nan 54. 63. 65. 58. nan nan
 59. nan]
[ True False False  True  True False  True False False  True False False
  True  True  True  True False False  True False]
```

```
[87]: b = np.where(a < 50, True, False)
      print(a)
      print(b)
```

```
[78 46 26 63 86  2 96 45 13 67 37 36 54 63 65 58 49 48 59 26]
[False  True  True False False  True False  True  True False  True  True]
```



```
False False False False  True  True False  True]
```

```
[88]: b = np.where(a < 50, 0, 100)
      print(a)
      print(b)
```

```
[78 46 26 63 86 2 96 45 13 67 37 36 54 63 65 58 49 48 59 26]
[100 0 0 100 100 0 100 0 0 100 0 0 100 100 100 100 0 0
 100 0]
```

1.0.8 Some operations with arrays

```
[90]: a
```

```
[90]: array([78, 46, 26, 63, 86,  2, 96, 45, 13, 67, 37, 36, 54, 63, 65, 58, 49,
            48, 59, 26])
```

```
[91]: a + 1
```

```
[91]: array([79, 47, 27, 64, 87,  3, 97, 46, 14, 68, 38, 37, 55, 64, 66, 59, 50,
            49, 60, 27])
```

```
[92]: a**2 + 3*a**3
```

```
[92]: array([1429740, 294124, 53404, 754110, 1915564, 28, 2663424,
           275400, 6760, 906778, 153328, 141264, 475308, 754110,
           828100, 588700, 355348, 334080, 619618, 53404])
```

```
[93]: # look for the integers I so that  $i^2 + (i+1)^2 = (i+2)^2$ 
      i = np.arange(30)
      b =  $i^2 + (i+1)^2$ 
```

```
[94]: c = (i+2)**2
```

```
[95]: print(b)
      print(c)
```

```
[ 1 5 13 25 41 61 85 113 145 181 221 265 313 365
 421 481 545 613 685 761 841 925 1013 1105 1201 1301 1405 1513
1625 1741]
[ 4 9 16 25 36 49 64 81 100 121 144 169 196 225 256 289 324 361
400 441 484 529 576 625 676 729 784 841 900 961]
```

```
[96] : b == c
```

```
[96]: array([False, False, False,  True, False, False, False, False, False,
          False, False, False, False, False, False, False, False, False,
          False, False, False, False, False, False, False, False, False,
```

```
False, False, False])
```

```
[97]: i[b==c]
```

```
[97]: array([3])
```

```
[98]: i[b==c][0] # the result is an array. To obtain the first value (here the only  
↪one), use [0]
```

```
[98]: 3
```

Numpy manages almost any mathematical operation. log, trigo, etc

```
[99]: a = np.arange(18)  
print(a)  
print(np.log10(a))
```

```
[ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17]  
[      -inf  0.          0.30103    0.47712125  0.60205999  0.69897  
 0.77815125  0.84509804  0.90308999  0.95424251  1.          1.04139269  
 1.07918125  1.11394335  1.14612804  1.17609126  1.20411998  1.23044892]
```

/Users/christophemorriset/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:3: RuntimeWarning: divide by zero encountered in log10

This is separate from the ipykernel package so we can avoid doing imports until

```
[100]: for aa in a:  
        print('{0:2} {1:4.2f} {2:5.2f} {3:8.2e}'.format(aa, np.log10(aa), np.  
↪sin(aa), np.exp(aa)))
```

```
0 -inf  0.00 1.00e+00  
1 0.00  0.84 2.72e+00  
2 0.30  0.91 7.39e+00  
3 0.48  0.14 2.01e+01  
4 0.60 -0.76 5.46e+01  
5 0.70 -0.96 1.48e+02  
6 0.78 -0.28 4.03e+02  
7 0.85  0.66 1.10e+03  
8 0.90  0.99 2.98e+03  
9 0.95  0.41 8.10e+03  
10 1.00 -0.54 2.20e+04  
11 1.04 -1.00 5.99e+04  
12 1.08 -0.54 1.63e+05  
13 1.11  0.42 4.42e+05  
14 1.15  0.99 1.20e+06  
15 1.18  0.65 3.27e+06
```

```
16 1.20 -0.29 8.89e+06
17 1.23 -0.96 2.42e+07
```

```
/Users/christophemorisset/anaconda3/lib/python3.7/site-
packages/ipykernel_launcher.py:2: RuntimeWarning: divide by zero encountered in
log10
```

```
sum
```

```
[103]: print(a.sum())
       print(17*18/2)
```

```
153
153.0
```

```
[104]: a = np.random.rand(2, 4, 3)
       print(a.shape)
       print(a.size)
```

```
(2, 4, 3)
24
```

```
2 planes, 4 rows, 3 columns
```

A small comment on the order of the elements in arrays in Python: There is two ways arrays can be stored: row- or column major. It has a direct impact on the way one has to loop on the arrays. IDL is like Fortran (column major) and Python is like C (row major). It means that in Python, as you move linearly through the memory of an array, the second dimension (rightmost) changes the fastest, while in IDL the first (leftmost) dimension changes the fastest. Consequence on the loop order in Python:

```
[105]: for plane in a:
       for row in plane:
           for col in row:
               print(col)
               print('-----')
```

```
0.9005117731906158
-----
0.318277757182466
-----
0.17908201684049108
-----
0.45874134723090876
-----
0.6789809335383882
-----
0.414421162986482
-----
```

```

0.056519355091393386
-----
0.8540270110679247
-----
0.8856219085188352
-----
0.888237760139903
-----
0.4771966465839834
-----
0.7799670136275508
-----
0.39921932133177807
-----
0.266255324504438
-----
0.27108565187535727
-----
0.6471891787778047
-----
0.04025540391379612
-----
0.38441391453496976
-----
0.6829276670149604
-----
0.6936363096727493
-----
0.8700864875082074
-----
0.7815463993013263
-----
0.17645130549021926
-----
0.240524011415837
-----

```

```
[ ]: print(a[0,1,2]) # a[p, r, c]
```

```
[106]: a.sum()
```

```
[106]: 12.345175661340384
```

```
[107]: a.sum(0) # from 3D to 2D. Generate an "image" of the sum, i.e. the "projection"
        ↪ on the x-axis of the 3D array
```

```
[107]: array([[1.29973109, 0.58453308, 0.45016767],
             [1.10593053, 0.71923634, 0.79883508],
             [0.73944702, 1.54766332, 1.7557084 ],
             [1.66978416, 0.65364795, 1.02049103]])
```

```
[108]: a.sum(0).shape
```

```
[108]: (4, 3)
```

```
[109]: a.sum(0).sum(0) # from 3D to 1D. From the image, make the sum in each row.
```

```
[109]: array([4.8148928 , 3.50508069, 4.02520217])
```

```
[110]: a.min(0)
```

```
[110]: array([[0.39921932, 0.26625532, 0.17908202],
             [0.45874135, 0.0402554 , 0.38441391],
             [0.05651936, 0.69363631, 0.87008649],
             [0.7815464 , 0.17645131, 0.24052401]])
```

```
[111]: a.ravel()
```

```
[111]: array([0.90051177, 0.31827776, 0.17908202, 0.45874135, 0.67898093,
             0.41442116, 0.05651936, 0.85402701, 0.88562191, 0.88823776,
             0.47719665, 0.77996701, 0.39921932, 0.26625532, 0.27108565,
             0.64718918, 0.0402554 , 0.38441391, 0.68292767, 0.69363631,
             0.87008649, 0.7815464 , 0.17645131, 0.24052401])
```

```
[112]: i_min = a.argmin() # return the index of where the minimum is. It uses the 1D
      ↪ index.
      print(i_min)
      b = np.array([10,2,3,4,5,2])
      b.argmin() # only the first occurrence
```

```
16
```

```
[112]: 1
```

```
[113]: a.ravel().shape # 1D
```

```
[113]: (24,)
```

```
[114]: a.ravel()[i_min] # Check where the minimum is.
```

```
[114]: 0.04025540391379612
```

```
[115]: z = i_min // 12
        y = (i_min - 12*z) // 3
        x = i_min - 12*z - 3*y
        print(z, y, x)
        print(a[z, y, x])
```

```
1 1 1
0.04025540391379612
```

```
[116]: def decompose_ravel(arr, i):
        shapes = arr.shape
        idx = i
        res = []
        for i in np.arange(arr.ndim):
            subdims = np.prod(shapes[i+1:])
            n = int(idx // subdims)
            #print n, subdims, idx
            idx = idx - subdims*n
            res.append(n)
        return tuple(res)
```

```
[117]: res = decompose_ravel(a, i_min)
        print(a.min())
        print(res)
        print(a[res])
```

```
0.04025540391379612
(1, 1, 1)
0.04025540391379612
```

```
[118]: a.min(0).min(0)
```

```
[118]: array([0.05651936, 0.0402554 , 0.17908202])
```

```
[119]: print(a[:,0,0])
        a[:,0,0].min()
```

```
[0.90051177 0.39921932]
```

```
[119]: 0.39921932133177807
```

```
[120]: a.mean(0)
```

```
[120]: array([[0.64986555, 0.29226654, 0.22508383],
              [0.55296526, 0.35961817, 0.39941754],
              [0.36972351, 0.77383166, 0.8778542 ],
              [0.83489208, 0.32682398, 0.51024551]])
```

```
[121]: np.median(a, 1)
```

```
[121]: array([[0.67348955, 0.57808879, 0.59719409],  
            [0.66505842, 0.22135331, 0.32774978]])
```

```
[124]: a.std(2)
```

```
[124]: array([[0.31248694, 0.11569194, 0.38361261, 0.17395633],  
            [0.0615729 , 0.24852108, 0.08581489, 0.27140593]])
```

```
[125]: np.percentile(a, 25)
```

```
[125]: 0.2698780700326274
```

```
[126]: print(a[0:4,0])  
print(np.cumsum(a[0:100,0])) # axis is a keyword. If absent, applied on the  
    ↪ ravel(), e.g. 1D array
```

```
[[0.90051177 0.31827776 0.17908202]  
 [0.39921932 0.26625532 0.27108565]]  
[0.90051177 1.21878953 1.39787155 1.79709087 2.06334619 2.33443184]
```

```
[127]: b = np.arange(1000).reshape(10,10,10)
```

```
[128]: b.shape
```

```
[128]: (10, 10, 10)
```

```
[129]: b[4,:,:] # hundreds digits = 4
```

```
[129]: array([[400, 401, 402, 403, 404, 405, 406, 407, 408, 409],  
            [410, 411, 412, 413, 414, 415, 416, 417, 418, 419],  
            [420, 421, 422, 423, 424, 425, 426, 427, 428, 429],  
            [430, 431, 432, 433, 434, 435, 436, 437, 438, 439],  
            [440, 441, 442, 443, 444, 445, 446, 447, 448, 449],  
            [450, 451, 452, 453, 454, 455, 456, 457, 458, 459],  
            [460, 461, 462, 463, 464, 465, 466, 467, 468, 469],  
            [470, 471, 472, 473, 474, 475, 476, 477, 478, 479],  
            [480, 481, 482, 483, 484, 485, 486, 487, 488, 489],  
            [490, 491, 492, 493, 494, 495, 496, 497, 498, 499]])
```

```
[130]: b[:,2,:] # tens digit = 2
```

```
[130]: array([[ 20,  21,  22,  23,  24,  25,  26,  27,  28,  29],  
            [120, 121, 122, 123, 124, 125, 126, 127, 128, 129],  
            [220, 221, 222, 223, 224, 225, 226, 227, 228, 229],  
            [320, 321, 322, 323, 324, 325, 326, 327, 328, 329],
```

```
[420, 421, 422, 423, 424, 425, 426, 427, 428, 429],
[520, 521, 522, 523, 524, 525, 526, 527, 528, 529],
[620, 621, 622, 623, 624, 625, 626, 627, 628, 629],
[720, 721, 722, 723, 724, 725, 726, 727, 728, 729],
[820, 821, 822, 823, 824, 825, 826, 827, 828, 829],
[920, 921, 922, 923, 924, 925, 926, 927, 928, 929]])
```

```
[131]: b[:, :, 7] # unity digit = 7
```

```
[131]: array([[ 7, 17, 27, 37, 47, 57, 67, 77, 87, 97],
 [107, 117, 127, 137, 147, 157, 167, 177, 187, 197],
 [207, 217, 227, 237, 247, 257, 267, 277, 287, 297],
 [307, 317, 327, 337, 347, 357, 367, 377, 387, 397],
 [407, 417, 427, 437, 447, 457, 467, 477, 487, 497],
 [507, 517, 527, 537, 547, 557, 567, 577, 587, 597],
 [607, 617, 627, 637, 647, 657, 667, 677, 687, 697],
 [707, 717, 727, 737, 747, 757, 767, 777, 787, 797],
 [807, 817, 827, 837, 847, 857, 867, 877, 887, 897],
 [907, 917, 927, 937, 947, 957, 967, 977, 987, 997]])
```

```
[132]: b.min(0) # elements with the smallest value for the hundreds digit
```

```
[132]: array([[ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9],
 [10, 11, 12, 13, 14, 15, 16, 17, 18, 19],
 [20, 21, 22, 23, 24, 25, 26, 27, 28, 29],
 [30, 31, 32, 33, 34, 35, 36, 37, 38, 39],
 [40, 41, 42, 43, 44, 45, 46, 47, 48, 49],
 [50, 51, 52, 53, 54, 55, 56, 57, 58, 59],
 [60, 61, 62, 63, 64, 65, 66, 67, 68, 69],
 [70, 71, 72, 73, 74, 75, 76, 77, 78, 79],
 [80, 81, 82, 83, 84, 85, 86, 87, 88, 89],
 [90, 91, 92, 93, 94, 95, 96, 97, 98, 99]])
```

```
[133]: b.min(2) # smallest value for the unity digit
```

```
[133]: array([[ 0, 10, 20, 30, 40, 50, 60, 70, 80, 90],
 [100, 110, 120, 130, 140, 150, 160, 170, 180, 190],
 [200, 210, 220, 230, 240, 250, 260, 270, 280, 290],
 [300, 310, 320, 330, 340, 350, 360, 370, 380, 390],
 [400, 410, 420, 430, 440, 450, 460, 470, 480, 490],
 [500, 510, 520, 530, 540, 550, 560, 570, 580, 590],
 [600, 610, 620, 630, 640, 650, 660, 670, 680, 690],
 [700, 710, 720, 730, 740, 750, 760, 770, 780, 790],
 [800, 810, 820, 830, 840, 850, 860, 870, 880, 890],
 [900, 910, 920, 930, 940, 950, 960, 970, 980, 990]])
```

```
[134]: b.min(2).shape
```



```
[134]: (10, 10)
```

```
[135]: np.median(b)
```

```
[135]: 499.5
```

```
[136]: np.median(b, axis=0)
```

```
[136]: array([[450., 451., 452., 453., 454., 455., 456., 457., 458., 459.],
        [460., 461., 462., 463., 464., 465., 466., 467., 468., 469.],
        [470., 471., 472., 473., 474., 475., 476., 477., 478., 479.],
        [480., 481., 482., 483., 484., 485., 486., 487., 488., 489.],
        [490., 491., 492., 493., 494., 495., 496., 497., 498., 499.],
        [500., 501., 502., 503., 504., 505., 506., 507., 508., 509.],
        [510., 511., 512., 513., 514., 515., 516., 517., 518., 519.],
        [520., 521., 522., 523., 524., 525., 526., 527., 528., 529.],
        [530., 531., 532., 533., 534., 535., 536., 537., 538., 539.],
        [540., 541., 542., 543., 544., 545., 546., 547., 548., 549.]])
```

```
[137]: x = 2 * np.random.rand(100,100,100) - 1.
      print(np.min(x), np.max(x))
```

```
-0.9999993984462574 0.9999991112403628
```

```
[138]: y = 2 * np.random.rand(100,100,100) - 1.
      z = 2 * np.random.rand(100,100,100) - 1.
```

```
[139]: r = np.sqrt(x**2 + y**2 + z**2)
      print(np.min(r), np.max(r))
      print(np.sqrt(3))
```

```
0.01262140022137408 1.7179592480943136
1.7320508075688772
```

```
[140]: print(np.mean(r))
      print(r.mean())
```

```
0.9607950001733658
0.9607950001733658
```

```
[141]: np.median(r)
```

```
[141]: 0.9848986271431883
```

1.0.9 Broadcasting

<http://arxiv.org/pdf/1102.1523.pdf>

If the two arrays differ in their number of dimensions, the shape of the array with fewer dimensions is prepended with ones until it matches the number of dimensions of the other array. If the shape of the two arrays does not match in any dimension, the array with shape equal to 1 in that dimension is prepended with ones until it matches the number of dimensions of the other array. If in any dimension the sizes disagree and neither is equal to 1, an error is raised.

```
[142]: x1 = np.array((1,2,3,4,5))
      y1 = np.array((1,2,3,4,5))
      z1 = np.array((1,2,3,4,5))
      r1 = x1 * y1 * z1
      print(r1.shape)
```

```
(5,)
```

```
[143]: x = np.array((1,2,3,4,5)).reshape(5,1,1)
```

```
[144]: x
```

```
[144]: array([[[1]],
              [[2]],
              [[3]],
              [[4]],
              [[5]])])
```

```
[145]: x.shape
```

```
[145]: (5, 1, 1)
```

```
[146]: x.ndim
```

```
[146]: 3
```

```
[147]: y = np.array((1,2,3,4,5)).reshape(1,5,1)
      z = np.array((1,2,3,4,5)).reshape(1,1,5)
      print(y)
      print(z)
```

```
[[[1]
  [2]
  [3]
  [4]
  [5]]]
[[[1 2 3 4 5]]]
```

```
[148]: r = x * y * z
```

```
[149]: print(r.shape)
```

```
(5, 5, 5)
```

```
[150]: r
```

```
[150]: array([[[ 1,  2,  3,  4,  5],
               [ 2,  4,  6,  8, 10],
               [ 3,  6,  9, 12, 15],
               [ 4,  8, 12, 16, 20],
               [ 5, 10, 15, 20, 25]],

              [[ 2,  4,  6,  8, 10],
               [ 4,  8, 12, 16, 20],
               [ 6, 12, 18, 24, 30],
               [ 8, 16, 24, 32, 40],
               [10, 20, 30, 40, 50]],

              [[ 3,  6,  9, 12, 15],
               [ 6, 12, 18, 24, 30],
               [ 9, 18, 27, 36, 45],
               [12, 24, 36, 48, 60],
               [15, 30, 45, 60, 75]],

              [[ 4,  8, 12, 16, 20],
               [ 8, 16, 24, 32, 40],
               [12, 24, 36, 48, 60],
               [16, 32, 48, 64, 80],
               [20, 40, 60, 80, 100]],

              [[ 5, 10, 15, 20, 25],
               [10, 20, 30, 40, 50],
               [15, 30, 45, 60, 75],
               [20, 40, 60, 80, 100],
               [25, 50, 75, 100, 125]]])
```

```
[151]: a = np.ones((10,10))
       b = np.arange(10).reshape(10,1)
       print(a)
       print(b)
       print(b.shape)
```

```
[[1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]
 [1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]
 [1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]
 [1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]
 [1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]
 [1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]
```

```

[1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]
[1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]
[1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]
[1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]]
[[0]
 [1]
 [2]
 [3]
 [4]
 [5]
 [6]
 [7]
 [8]
 [9]]
(10, 1)

```

```
[153]: a * b
```

```

[153]: array([[0., 0., 0., 0., 0., 0., 0., 0., 0., 0.],
 [1., 1., 1., 1., 1., 1., 1., 1., 1., 1.],
 [2., 2., 2., 2., 2., 2., 2., 2., 2., 2.],
 [3., 3., 3., 3., 3., 3., 3., 3., 3., 3.],
 [4., 4., 4., 4., 4., 4., 4., 4., 4., 4.],
 [5., 5., 5., 5., 5., 5., 5., 5., 5., 5.],
 [6., 6., 6., 6., 6., 6., 6., 6., 6., 6.],
 [7., 7., 7., 7., 7., 7., 7., 7., 7., 7.],
 [8., 8., 8., 8., 8., 8., 8., 8., 8., 8.],
 [9., 9., 9., 9., 9., 9., 9., 9., 9., 9.]])

```

```
[154]: a * b.reshape(1,10)
```

```

[154]: array([[0., 1., 2., 3., 4., 5., 6., 7., 8., 9.],
 [0., 1., 2., 3., 4., 5., 6., 7., 8., 9.],
 [0., 1., 2., 3., 4., 5., 6., 7., 8., 9.],
 [0., 1., 2., 3., 4., 5., 6., 7., 8., 9.],
 [0., 1., 2., 3., 4., 5., 6., 7., 8., 9.],
 [0., 1., 2., 3., 4., 5., 6., 7., 8., 9.],
 [0., 1., 2., 3., 4., 5., 6., 7., 8., 9.],
 [0., 1., 2., 3., 4., 5., 6., 7., 8., 9.],
 [0., 1., 2., 3., 4., 5., 6., 7., 8., 9.],
 [0., 1., 2., 3., 4., 5., 6., 7., 8., 9.]])

```

1.0.10 Structured arrays and RecArrays

See here: <http://docs.scipy.org/doc/numpy/user/basics.rec.html>

A structured array in numpy is an array of records. Each record can contain one or more items which can be of different types.

```
[155]: a = np.array([(1.5, 2), (3.0, 4)]) # Classical numpy array
print(a)
```

```
[[1.5 2. ]
 [3.  4. ]]
```

```
[156]: astru = np.array([(1.5, 2), (3.0, 4)], dtype=[('x', float), ('y', int)]) #
↳ array with named and typed columns
astru
```

```
[156]: array([(1.5, 2), (3. , 4)], dtype=[('x', '<f8'), ('y', '<i8')])
```

```
[157]: print(astru['x'])
print(astru['y'])
```

```
[1.5 3. ]
[2 4]
```

```
[158]: arec = astru.view(np.recarray)
print(type(a), type(astru), type(arec))
print('-----')
print(a)
print(astru)
print(arec)
print('-----')
print(a.size, astru.size, arec.size) # not even the same size
print('-----')
print(a.dtype, astru.dtype, arec.dtype) # types tell us that ar has column
↳ names and types
print('-----')
print(a[1,1], astru[1][1], arec[1][1]) # one is 2D, the other is a collection
↳ of 1D
print('-----')
print(astru['y']) # acces by name (a little like dictionnaires)
print('-----')
print(arec.x)
```

```
<class 'numpy.ndarray'> <class 'numpy.ndarray'> <class 'numpy.recarray'>
```

```
-----
[[1.5 2. ]
 [3.  4. ]]
[(1.5, 2) (3. , 4)]
[(1.5, 2) (3. , 4)]
-----
```

```
4 2 2
-----
```

```
float64 [('x', '<f8'), ('y', '<i8')] (numpy.record, [('x', '<f8'), ('y',
'<i8')])
```

```
-----  
4.0 4 4  
-----
```

```
[2 4]  
-----
```

```
[1.5 3. ]
```

```
[ ]: %timeit astru2 = np.append(astru, np.array([(5.0, 6)], dtype=astru.dtype)) #  
      ↳ Copied all the data, may be slow
```

```
[ ]: %timeit astru3 = np.concatenate((astru, np.array([(5.0, 6)], dtype=astru.  
      ↳ dtype))) # A little bit faster
```

```
[ ]: %timeit arec2 = np.append(arec, np.array([(5.0, 6)], dtype=astru.dtype).view(np.  
      ↳ recarray)) # Copied all the data, may be slow
```

```
[ ]: %timeit arec3 = np.concatenate((arec, np.array([(5.0, 6)], dtype=astru.dtype).  
      ↳ view(np.recarray))) # A little bit faster
```

```
[159]: arec4 = np.rec.fromrecords([(456, 'dbe', 1.2), (2, 'de', 1.  
      ↳ 3)], names='col1,col2,col3') # direct from data.  
print(arec4)  
print(type(arec4))  
print(arec4.col1[1])  
print(arec4[1].col1)
```

```
[(456, 'dbe', 1.2) ( 2, 'de', 1.3)]
```

```
<class 'numpy.recarray'>
```

```
2
```

```
2
```

```
[160]: arec4 = np.rec.fromrecords([('etoile_15', 30.015, -0.752, 10.722),  
      ('etoile_11', 31.163, -9.109, 10.761),  
      ('etoile_16', 39.789, -7.716, 11.071),  
      ('etoile_14', 35.110, 6.785, 11.176),  
      ('etoile_31', 33.530, 9.306, 11.823),  
      ('etoile_04', 33.480, 5.568, 11.978)  
      ],  
      names='name,ra,dec, mag')
```

```
[ ]: mask = arec4.mag > 11.  
print(arec4[mask])  
print('-----')  
for star in arec4[mask]:  
    print('name: {0} ra = {1} dec = {2} magnitude = {3}'.format(star.name, star.  
      ↳ ra, star.dec, star.mag))  
print('-----')
```

```
for star in arec4[mask]:
    print('name: {0[name]} ra = {0[ra]} dec = {0[dec]} magnitude = {0[mag]}'.
    ↪format(star)) # unse only one key in format
```

1.0.11 NaN and other ANSI values

```
[161]: a = np.array([-3, -2., -1., 0., 1., 2.])
      b = 1./a
      print(b)
```

```
[-0.33333333 -0.5          -1.          inf  1.          0.5          ]

/Users/christophemorisset/anaconda3/lib/python3.7/site-
packages/ipykernel_launcher.py:2: RuntimeWarning: divide by zero encountered in
true_divide
```

```
[162]: print(a.sum())
      print(b.sum()) # NaN and others are absorbant elements
```

```
-3.0
inf
```

```
[163]: mask = np.isfinite(b)
      print(mask)
      print(b[mask].sum())
```

```
[ True  True  True False  True  True]
-0.33333333333333326
```

```
[164]: for elem in b:
      print(np.isinf(elem))
```

```
False
False
False
True
False
False
```

```
[165]: a = np.array([-2, -1, 1., 2, 3])
      b = np.log10(a)
      mask = np.isfinite(b)
      print(a)
      print(b)
      print(mask)
      print(a.mean())
      print(b.mean())
```

```
print(b[mask].mean())
print(np.nanmean(b))
```

```
[-2. -1.  1.  2.  3.]
[      nan      nan  0.      0.30103      0.47712125]
[False False  True  True  True]
0.6
nan
0.25938375012788123
0.25938375012788123

/Users/christophemorriset/anaconda3/lib/python3.7/site-
packages/ipykernel_launcher.py:2: RuntimeWarning: invalid value encountered in
log10
```

1.0.12 Roundish values of floats

```
[166]: import math
res = []
for i in range(100):
    res.append(math.log(2 ** i, 2)) # The second argument is the base of the
    ↪ log. The result should be i.
print(res)
# We can see that sometimes the value of log2(2**i) is NOT i.
```

```
[0.0, 1.0, 2.0, 3.0, 4.0, 5.0, 6.0, 7.0, 8.0, 9.0, 10.0, 11.0, 12.0, 13.0, 14.0,
15.0, 16.0, 17.0, 18.0, 19.0, 20.0, 21.0, 22.0, 23.0, 24.0, 25.0, 26.0, 27.0,
28.0, 29.000000000000004, 30.0, 31.000000000000004, 32.0, 33.0, 34.0, 35.0,
36.0, 37.0, 38.0, 39.000000000000001, 40.0, 41.0, 42.0, 43.0, 44.0, 45.0, 46.0,
47.000000000000001, 48.0, 49.0, 50.0, 51.000000000000001, 52.0, 53.0, 54.0,
55.000000000000001, 56.0, 57.0, 58.000000000000001, 59.000000000000001, 60.0, 61.0,
62.000000000000001, 63.0, 64.0, 65.0, 66.0, 67.0, 68.0, 69.0, 70.0, 71.0, 72.0,
73.0, 74.0, 75.0, 76.0, 77.0, 78.000000000000001, 79.0, 80.0, 81.0, 82.0, 83.0,
84.0, 85.0, 86.0, 87.0, 88.0, 89.0, 90.0, 91.0, 92.0, 93.000000000000001,
94.000000000000001, 95.000000000000001, 96.0, 97.0, 98.0, 99.0]
```

```
[167]: res2 = []
for i in range(100):
    res2.append(float(round(math.log(2**i, 2))) == math.log(2 ** i, 2))
print(res2)
# An equivalent result is obtained when comparing the round value. This should
    ↪ be always True.
```

```
[True, True, True, True, True, True, True, True, True, True, True, True, True,
True, True, True, True, True, True, True, True, True, True, True, True, True,
True, True, True, False, True, False, True, True, True, True, True, True, True,
False, True, True, True, True, True, True, True, False, True, True, True, False,
```


True, True, True, False, True, True, False, False, True, True, False, True,
True, True, True, True, True, True, True, True, True, True, True, True, True,
True, False, True, True, True, True, True, True, True, True, True, True, True,
True, True, True, False, False, False, True, True, True, True]

```
[168]: res = []
for i in range(100):
    res.append(np.log2(2.**i)) # The second argument is the base of the log.
    ↪ The result should be i.
print(res)

res_np = []
for i in range(100):
    res_np.append(float(round(np.log2(2.**i))) == np.log2(2.**i))
print(res_np)
# No problems with the numpy log function.
```

[illegible]

In case of doubts, one can use the `close` function from `numpy`:

```
[169]: res_np2 = []
        for i in range(100):
            res_np2.append(np.isclose(float(round(math.log(2 ** i, 2))), math.log(2 ** i, 2)))
        print(res_np2)
        # The isclose
```

[True, True, True, True, True, True, True, True, True, True, True, True, True,
True, True, True, True, True, True, True, True, True, True, True, True, True,
True, True, True, True, True, True, True, True, True, True, True, True, True,
True, True, True, True, True, True, True, True, True, True, True, True, True,
True, True, True, True, True, True, True, True, True, True, True, True, True,

```
True, True, True, True, True, True, True, True, True, True, True, True, True,
True, True, True, True, True, True, True, True, True, True, True, True, True,
True, True, True, True, True, True, True, True, True]
```

```
[170]: np.isclose?
```

Signature: `np.isclose(a, b, rtol=1e-05, atol=1e-08, equal_nan=False)`

Docstring:

Returns a boolean array where two arrays are element-wise equal within a tolerance.

The tolerance values are positive, typically very small numbers. The relative difference (``rtol` * abs(`b`)`) and the absolute difference ``atol`` are added together to compare against the absolute difference between ``a`` and ``b``.

.. warning:: The default ``atol`` is not appropriate for comparing numbers that are much smaller than one (see Notes).

Parameters

`a, b` : array_like

Input arrays to compare.

`rtol` : float

The relative tolerance parameter (see Notes).

`atol` : float

The absolute tolerance parameter (see Notes).

`equal_nan` : bool

Whether to compare NaN's as equal. If True, NaN's in ``a`` will be considered equal to NaN's in ``b`` in the output array.

Returns

`y` : array_like

Returns a boolean array of where ``a`` and ``b`` are equal within the given tolerance. If both ``a`` and ``b`` are scalars, returns a single boolean value.

See Also

`allclose`

Notes

.. versionadded:: 1.7.0

For finite values, `isclose` uses the following equation to test whether two floating point values are equivalent.

```
absolute(`a` - `b`) <= (`atol` + `rtol` * absolute(`b`))
```

Unlike the built-in `math.isclose`, the above equation is not symmetric in ``a`` and ``b`` -- it assumes ``b`` is the reference value -- so that `isclose(a, b)` might be different from `isclose(b, a)`. Furthermore, the default value of `atol` is not zero, and is used to determine what small values should be considered close to zero. The default value is appropriate for expected values of order unity: if the expected values are significantly smaller than one, it can result in false positives. `atol` should be carefully selected for the use case at hand. A zero value for `atol` will result in `False` if either ``a`` or ``b`` is zero.

Examples

```
-----
>>> np.isclose([1e10,1e-7], [1.00001e10,1e-8])
array([ True, False])
>>> np.isclose([1e10,1e-8], [1.00001e10,1e-9])
array([ True, True])
>>> np.isclose([1e10,1e-8], [1.0001e10,1e-9])
array([False,  True])
>>> np.isclose([1.0, np.nan], [1.0, np.nan])
array([ True, False])
>>> np.isclose([1.0, np.nan], [1.0, np.nan], equal_nan=True)
array([ True, True])
>>> np.isclose([1e-8, 1e-7], [0.0, 0.0])
array([ True, False])
>>> np.isclose([1e-100, 1e-7], [0.0, 0.0], atol=0.0)
array([False, False])
>>> np.isclose([1e-10, 1e-10], [1e-20, 0.0])
array([ True,  True])
>>> np.isclose([1e-10, 1e-10], [1e-20, 0.999999e-10], atol=0.0)
array([False,  True])
File:      ~/anaconda3/lib/python3.7/site-packages/numpy/core/numeric.py
Type:      function
```